



Smaller weights and faster execution are separate results

SMILES-only Mol-JEPA · float32 · all output heads retained

56.47% fewer parameters



Measured Apple MPS speedup



Inactive encoders save storage. Sparse graph operations and batched readouts provide additional execution gains.

Latency: median and p10-p90 over 20 interleaved rounds after 10 warmups, including parsing and transfers; Apple M5, 16 GB.

Source: benchmarks/moljepa_runtime_macos.json. Apple MPS only; NVIDIA CUDA unverified. No per-rewrite speed attribution is inferred.